

Kesava Sai Krishna Kondepudi

(+91) 8008407440 | kesavakrishna91@gmail.com | [linkedin.com/in/kesava-krishna](https://www.linkedin.com/in/kesava-krishna) | github.com/kesavakrishna

SUMMARY

AI Engineer with 2+ years building production LLM systems and 2 peer-reviewed publications. Built enterprise RAG platforms, multi-agent workflows, and geospatial intelligence systems used by 90+ users across 25+ clients. Interested in reliable AI systems that perform under distribution shift.

EDUCATION

Mahindra University

Bachelor of Technology in Electronics and Computer Engineering

Hyderabad, India

Oct 2020 – Jun 2024

EXPERIENCE

Bilvantis Technologies

Programmer Analyst: Artificial Intelligence

Hyderabad, India

Jun 2024 – Present

- Architected **Tiggo**, a location intelligence platform combining Google Maps API, environmental datasets, and demographic data to automate site feasibility analysis and revenue projection for enterprise clients.
- Designed Tiggo's fault-tolerant Temporal-orchestrated ETL pipelines for concurrent ingestion from 10+ third-party APIs and unstructured document sources, with retry semantics for transient failures.
- Deployed Tiggo's RFP intelligence module using Perplexity, Jina, and OpenAI-based relevance scoring to parse multi-format documents — automating end-to-end RFP processing for 90+ enterprise users and reducing manual review by 60%.
- Developed **P2X**, an enterprise migration platform using LLM-based code analysis and automated SQL generation to convert legacy Informatica workflows to cloud databases with 95% accuracy — reducing migration time by 70% and infrastructure costs by 60% across 25+ clients.
- Built **RADEX**, a multi-tenant conversational AI platform using partitioned Milvus vector stores, hybrid retrieval, and a Text-to-SQL agent to serve context-aware responses across structured and unstructured enterprise data; diagnosed production failures caused by distribution shift in client query patterns.
- Engineered an LLM-powered test automation system that interprets Jira/ClickUp tickets to generate pytest unit tests, integration tests, and Postman collections — reducing manual QA effort by 50%.

AI/ML Intern

Mar 2024 – Jun 2024

- Developed an ensemble anomaly detection system combining LSTMs, autoencoders, and tree-based models (XGBoost, Isolation Forest) for highly imbalanced time-series data, achieving 95% detection accuracy — later extended into the full-time role.

Feynn Labs

Machine Learning Intern

Remote

Jul 2023 – Sep 2023

- Built a market segmentation model for the Indian EV market using K-Means with PCA dimensionality reduction; identified 4 optimal customer segments via Elbow Method across range, price, charging access, and form factor — producing geographic and behavioral profiles to inform market entry strategy.

PUBLICATIONS

Shaik, M., Sahjwani, A., **Kondepudi, K.S.K.** Forecasting ASEAN-5 Stock Index Price Movement Using Machine Learning Techniques. *Journal of Prediction Markets*, 18(1), 2024. — Benchmarked 8 ML models across 11 years of Southeast Asian market data; AdaBoost achieved superior directional prediction accuracy. doi: [10.5750/jpm.v18i1.2119](https://doi.org/10.5750/jpm.v18i1.2119).

Sahjwani, A., Choudhary, N., **Kondepudi, K.S.K.**, et al. A Deep Learning Based Approach for Cardiovascular Monitoring Using Seismocardiogram Signals. *IEEE ICPEEV*, 2023. — CycleGAN-based pipeline translating non-invasive SCG signals to ECG representations, achieving 89% reconstruction fidelity (96.55% in extended work). doi: [10.1109/ICPEEV58650.2023.10391941](https://doi.org/10.1109/ICPEEV58650.2023.10391941).

PROJECTS

Reflective AI Memory System

Jun 2025 – Present

- Memory-augmented AI agent (Gemini + FAISS) with cross-session recall, semantic memory segmentation, and a personalized prompting strategy integrating past context. Designing next-phase self-modeling for autonomous learning.

Automated Corporate Risk Assessment

Apr 2025 – Jun 2025

- Multi-agent system (AutoGen) that analyzes financial indicators and scrapes news to identify market, operational, and reputational risks — generating an ERM-aligned risk matrix with mitigation strategies.

TECHNICAL SKILLS

Languages: Python, SQL, TypeScript

AI/ML: Large Language Models, RAG, Multi-Agent Systems, GANs, LSTMs, Autoencoders, NLP, Time-Series Forecasting, Ensemble Methods

Frameworks: PyTorch, TensorFlow, HuggingFace, Scikit-learn, LangChain, AutoGen, CrewAI

APIs & Tools: OpenAI, Gemini, Groq, Jina, AssemblyAI, Whisper, Git, Postman, Pandas, NumPy

Data & Infrastructure: Milvus, GCP (Compute Engine, Cloud Run, BigQuery, Vertex AI), Temporal, FastAPI, Docker